Docker Image Deployment Guide

Overview

This document provides step-by-step instructions for building, running, and deploying a Docker container locally and then pushing it to **AWS Elastic Container Registry (ECR)** for deployment on **AWS Elastic Container Service (ECS)**.

---

1. Build and Run Docker Image Locally

1.1 Build the Docker Image

To create a Docker image for your FastAPI application, run:

```
sudo docker build -t my-fastapi-app .
```

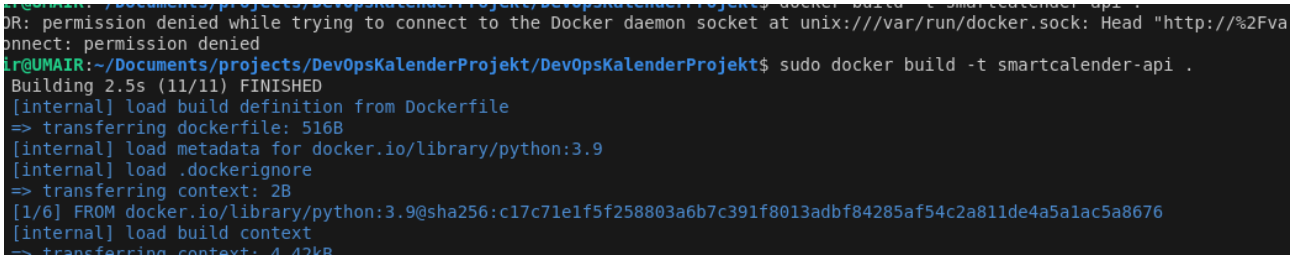

This command:

- Uses the `Dockerfile` in the current directory (`.`) to build the image.
- Tags the image as `my-fastapi-app`.

1.2 Run the Docker Container Locally

Execute the following command to run the container:

```
sudo docker run -p 8000:8000 my-fastapi-app
```

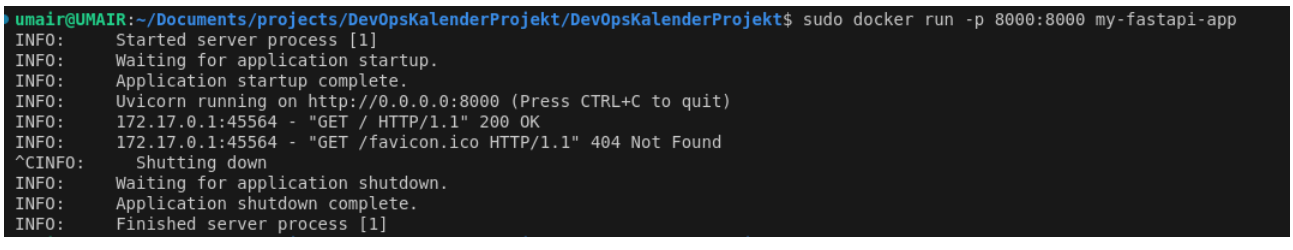

This maps:

- Port `8000` inside the container to port `8000` on the local machine.

The application should now be accessible at:

```
http://localhost:8000
```

---

2. Push Docker Image to AWS Elastic Container Registry (ECR)

2.1 Configure AWS CLI

Ensure AWS CLI is configured with valid credentials:

```
aws configure
```

This prompts for:

- AWS **Access Key**
- AWS **Secret Key**
- AWS **Region** (e.g., `eu-central-1`)

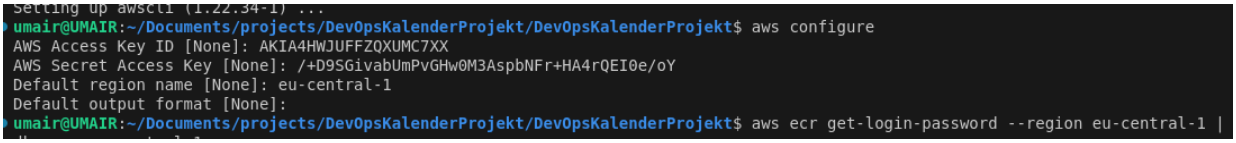

-

2.2 Authenticate Docker with AWS ECR

Use the following command to authenticate Docker with AWS ECR:

```
aws ecr get-login-password --region eu-central-1 | docker login --username AWS
--password-stdin 841162697075.dkr.ecr.eu-central-1.amazonaws.com
```

This logs Docker into AWS ECR to allow pushing images.

2.3 Build the Docker Image for AWS ECR

Run:

```
docker buildx build --platform linux/amd64 -t smartcalender-api .
```

This tags the image as `smartcalender-api`.

2.4 Tag the Docker Image for AWS ECR

```
docker tag smartcalender-api:latest
841162697075.dkr.ecr.eu-central-1.amazonaws.com/smartcalender-api:latest
```

This assigns the correct tag needed to push the image to ECR.

2.5 Push the Docker Image to AWS ECR

docker push 841162697075.dkr.ecr.eu-central-1.amazonaws.com/smartcalender-api:latest

```
Default output format [None]:
umair@UMAIR:~/Documents/projects/DevOpsKalenderProjekt/DevOpsKalenderProjekt$ aws ecr get-login-password --region eu-central-1 | docker login --username AWS --password-stdin 841162697075.
dkr.ecr.eu-central-1.amazonaws.com
WARNING! Your password will be stored unencrypted in /home/umair/.docker/config.json.
Configure a credential helper to remove this warning. See
https://docs.docker.com/engine/reference/commandline/login/#credential-stores

Login Succeeded
umair@UMAIR:~/Documents/projects/DevOpsKalenderProjekt/DevOpsKalenderProjekt$ docker build -t smartcalender-api .
ERROR: permission denied while trying to connect to the Docker daemon socket at unix:///var/run/docker.sock: Head "http://%2Fvar%2Frun%2Fdocker.sock/_ping": dial unix /var/run/docker.sock
: connect: permission denied
umair@UMAIR:~/Documents/projects/DevOpsKalenderProjekt/DevOpsKalenderProjekt$ sudo docker build -t smartcalender-api .
[+] Building 2.5s (11/11) FINISHED                                                                                                                    docker:default
 => [internal] load build definition from Dockerfile                                                                                                          0.0s
 => => transferring dockerfile: 516B                                                                                                                          0.0s
 => [internal] load metadata for docker.io/library/python:3.9                                                                                                 2.3s
 => [internal] load .dockerignore                                                                                                                             0.0s
 => => transferring context: 2B                                                                                                                               0.0s
 => [1/6] FROM docker.io/library/python:3.9@sha256:c17c71e1f5f258803a6b7c391f8013adbf84285af54c2a811de4a5a1ac5a8676                                           0.0s
 => [internal] load build context                                                                                                                             0.0s
 => => transferring context: 4.42kB                                                                                                                           0.0s
 => CACHED [2/6] WORKDIR /app                                                                                                                                 0.0s
 => CACHED [3/6] COPY requirements.txt .                                                                                                                      0.0s
 => CACHED [4/6] RUN pip install --no-cache-dir -r requirements.txt                                                                                           0.0s
 => CACHED [5/6] COPY . /app                                                                                                                                  0.0s
 => CACHED [6/6] COPY app/api_key.txt /app/app/api_key.txt                                                                                                    0.0s
 => exporting to image                                                                                                                                        0.0s
 => => exporting layers                                                                                                                                       0.0s
 => => writing image sha256:499e47f1199bf40c9e9d3c4a3187615c0d6a530c9244e8e4cd09fbc505b82f78                                                                  0.0s
 => => naming to docker.io/library/smartcalender-api                                                                                                          0.0s
umair@UMAIR:~/Documents/projects/DevOpsKalenderProjekt/DevOpsKalenderProjekt$ sudo docker tag smartcalender-api:latest 841162697075.dkr.ecr.eu-central-1.amazonaws.com/smartcalender-api:la
test
umair@UMAIR:~/Documents/projects/DevOpsKalenderProjekt/DevOpsKalenderProjekt$ sudo docker push 841162697075.dkr.ecr.eu-central-1.amazonaws.com/smartcalender-api:latest
The push refers to repository [841162697075.dkr.ecr.eu-central-1.amazonaws.com/smartcalender-api]
94bd550c68da: Preparing
59042177ee67: Preparing
c56a2f337329: Preparing
9c170551fc05: Preparing
7972b775c81c: Preparing
60a159600b22: Waiting
ee959616fc20: Waiting
d0e85779261a: Waiting
```

---

3. Deploy Docker Image to AWS ECS

After pushing the image to **AWS ECR**, use it in **AWS ECS** by:

1. Creating an **ECS Cluster**.

   Give name of cluster and click on create

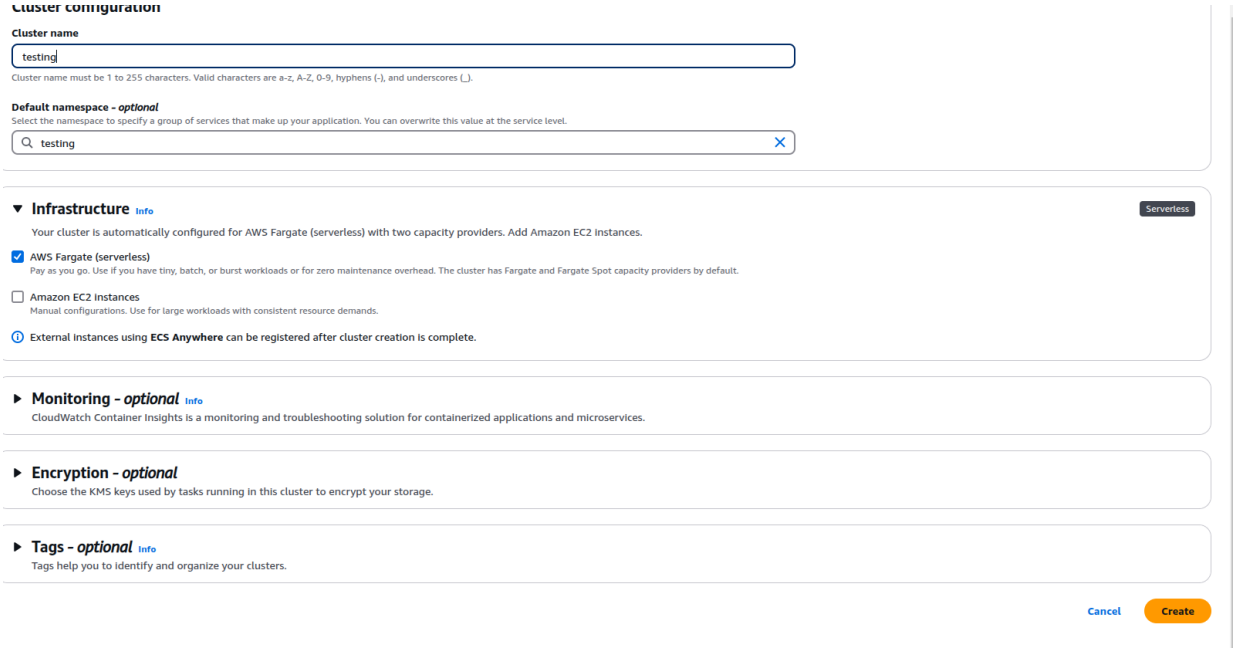

2. Defining a **Task Definition** referencing the ECR image.

   Give name of the **Task definition family**
   select **Task role** and **Task execution role** from dropdown

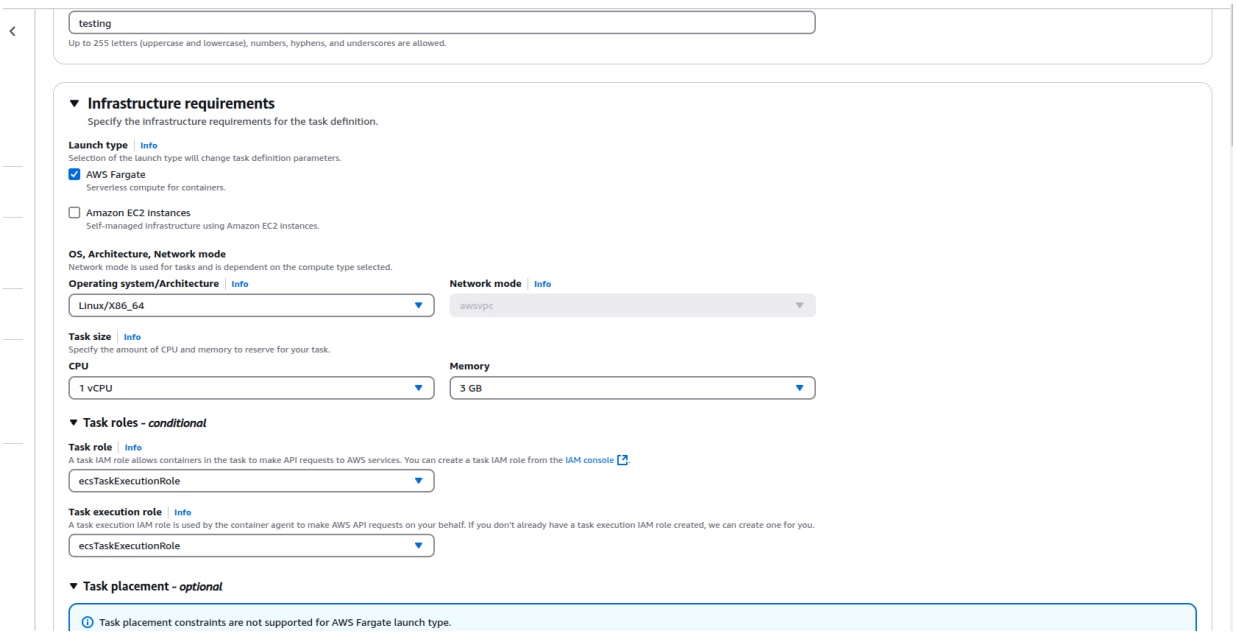

select the **CPU and Memory** according to your requirement

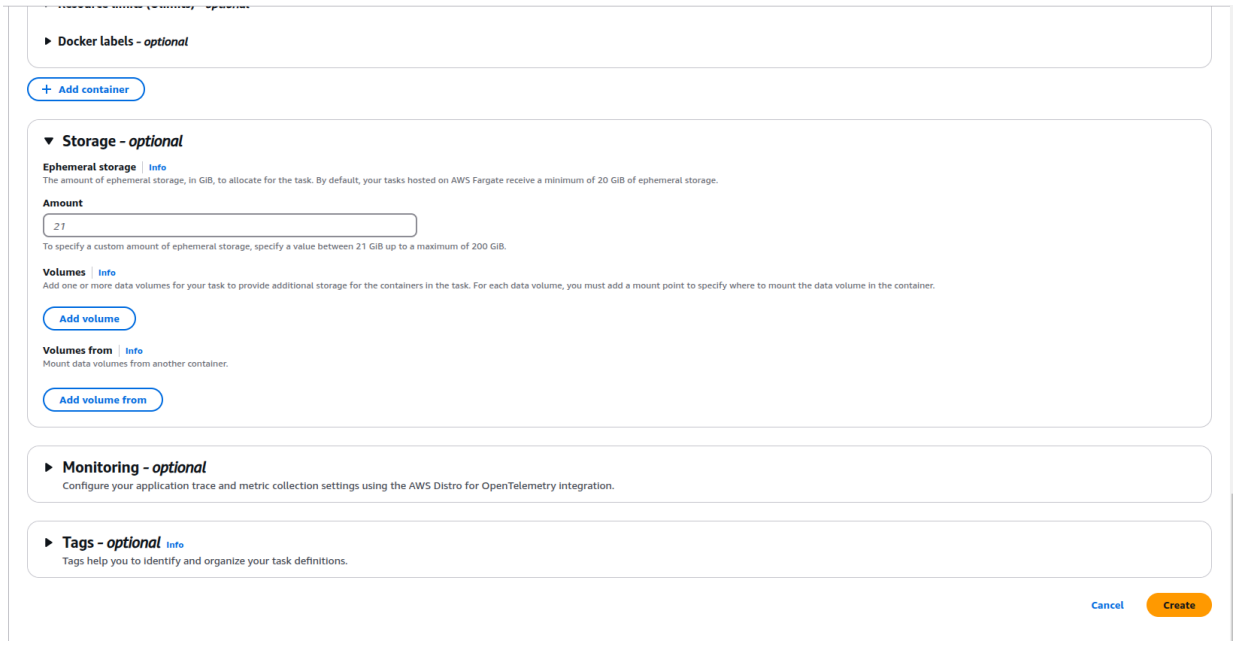

   leave everything same and click on done

3. Running an **ECS Service** using the Task Definition.

GO the cluster and and click on create to create service

select the task definition from **Family field** and select the **revision** from dropdown

give name of Service name in Service field

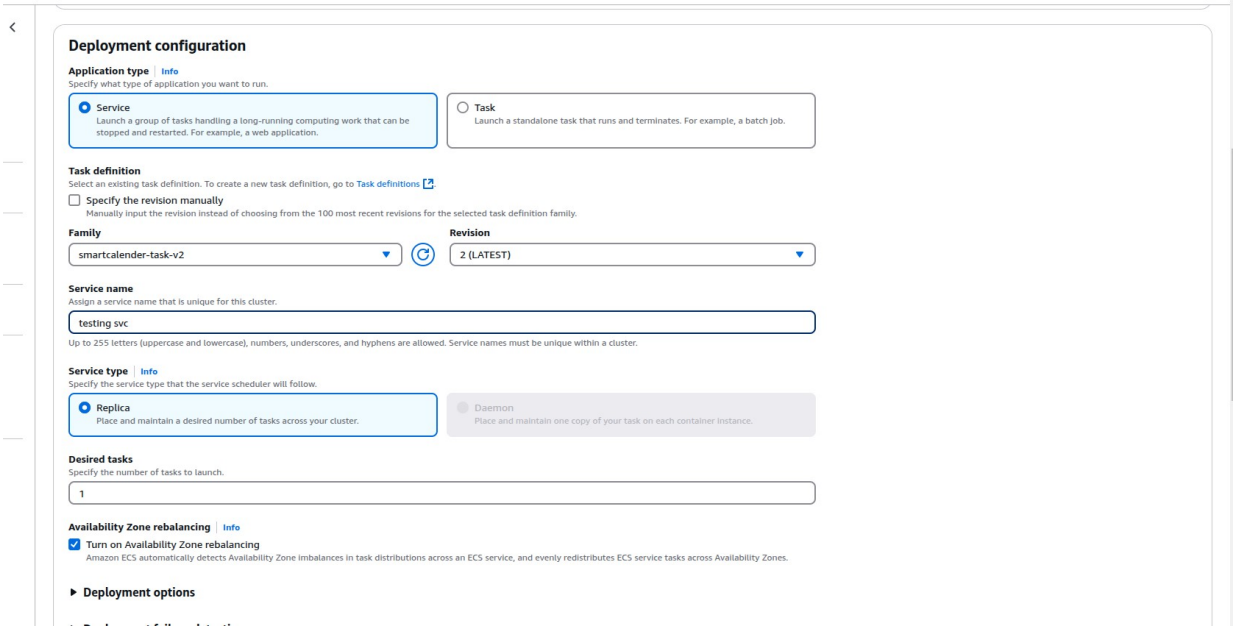

In Networking tab select the default vpc and subnet

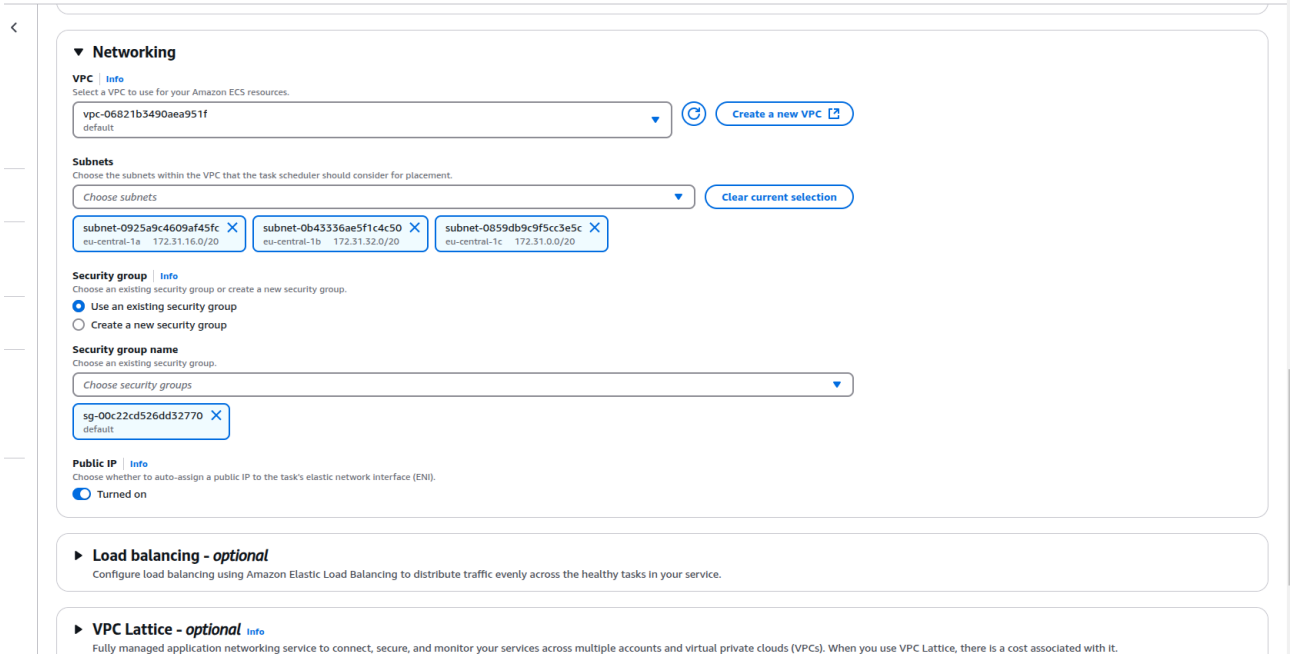

In load balancer section
give the name loadbalancer and leave other thing remain same

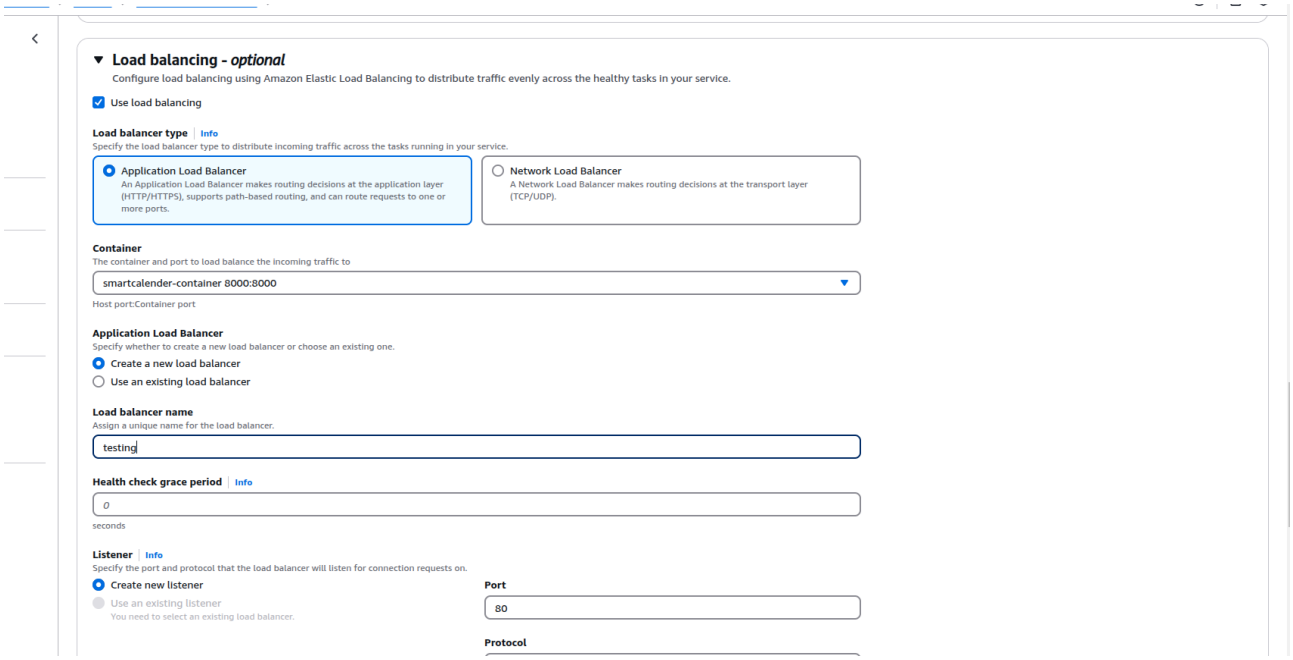

Leave everything same and click on create service

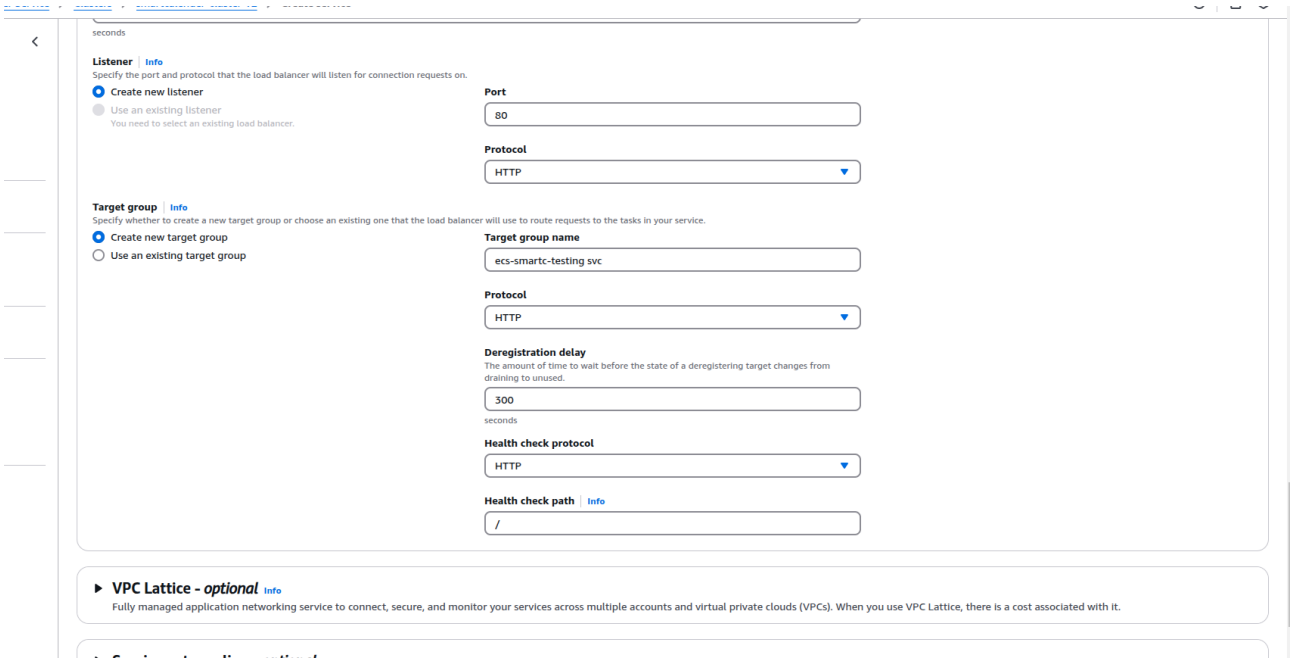

To copy the **Application Load Balancer (ALB) DNS name**, follow these steps based on your setup:

Go to the AWS Management Console → Navigate to **EC2**.

- In the left panel, click on **Load Balancers** (under Load Balancing).
- Locate your **Application Load Balancer (ALB)**.
- Click on the ALB name to open its details.
- Find the **DNS name** under the **Basic Configuration** section.
- Click the **copy icon** next to the DNS name to copy it.